

Improved running time for DP.

Add an extra step 3 : dominated pairs are removed. A pair (t, w) dominates another pair (t', w') if $t \leq t'$ and $w \geq w'$.

Theorem 3.1 The knapsack problem can be solved (exactly) in $O(n \cdot \min\{V, B\})$

Proof By removing each pair dominated by another pair, # of pairs in A_j is at most $\min\{B, V\}$. Thus step 1&2&3 can be done in $O(\min\{B, V\})$ time in each iteration.

Thus, the time is $O(n \cdot \min\{B, V\})$.

A PTAS for knapsack

1. Round the values $v'_i = \lfloor v_i / M \rfloor$, where $M = \frac{\epsilon M}{n}$, $M = \max_i v_i$.

2. Apply the DP to the second rounded instance.

Analysis

• After rounding, we have $v'_i \in \{0, 1, \dots, \lfloor n/\epsilon \rfloor\}$

→ DP runs in $O(nV') = O(n^3/\epsilon)$ time where $V' = \sum_i v'_i = O(n^2/\epsilon)$

→ Polynomial time $\hookrightarrow \min(B, V') \leq V'$.

• By rounding, the loss in precision is at most M per item.

→ loss is at most $nM = \epsilon \max\{v_i\} = \epsilon M \leq \epsilon \text{OPT}$.

→ $(1 - \epsilon)$ -approximation

Formal proof. Let S be the set of items found by DP and let O be an optimal set of items for the original (unrounded) instance.

For any item i , we have

$$\textcircled{1} \quad v_i \geq M \quad v'_i > M(v_i/M - 1) = v_i - M$$

The value of the final solution is

$$\sum_{i \in S} v_i \geq \sum_{i \in S} M v'_i \quad \textcircled{1} \quad \sum_{i \in O} M v'_i \geq \sum_{i \in O} (v_i - M)$$

$$\leq \text{OPT} - |O|M \geq \text{OPT} - nM = \text{OPT} - \epsilon M$$

$$\textcircled{2} \quad \geq \text{OPT} - \epsilon \text{OPT} = (1 - \epsilon) \text{OPT}$$

$\textcircled{1}$ S is optimal for the rounded instance with value v'_i

$\textcircled{2}$ $\text{OPT} \geq M$ since taking the item with the largest value is a feasible solution.

3.2 Scheduling on identical machines.

Define: job j is long if $p_j > ET^* \Rightarrow p_j > \epsilon P/m$

job j is short if $p_j \leq ET^* \Rightarrow p_j \leq \epsilon P/m$

choose T^* s.t. $T^* \leq \text{OPT}$, for example $T^* = P/m$ with $P = \sum_j p_j$.

PTAS: step 1: Find a schedule for the long jobs of length at most

$$(1 + \epsilon) \text{OPT} \quad C_{\max} \leq p_L + \sum_{j \neq L} p_j/m$$

step 2: Add short jobs greedily.

Analysis:

1. Feasible ✓

2. Ratio $\leq 1 + \epsilon$: the last job to complete is a short job.

$$\boxed{\sum_{j \neq L} p_j/m = S_L \leq P/m \leq \text{OPT}} \quad C_{\max} = S_L + p_L \leq S_L + ET^* \leq \text{OPT} + \epsilon \text{OPT} = (1 + \epsilon) \text{OPT}.$$

② the last job to complete is a long job.

the makespan equals the length of the optimal schedule for just the long jobs, which is clearly no more than C^* for entire input.

3. Running time: Lemma: An optimal schedule of the long jobs can be found in $O(m^{m/\epsilon})$ time.

Proof # of long jobs is at most $P/(ET^*) = m/\epsilon$.

To find the optimal one, try all possible schedules and take the best. Note the orders of jobs on a machine is not important for the length of the schedule. Thus, # of possible schedules is at most $m^{m/\epsilon}$.

Theorem 3.2 If m is assumed a constant, the PTAS runs in polynomial time.

The running time of the PTAS above is exponential in m . Next, we show how to get a PTAS with running time polynomial in m .

Ingredients:

- Partition in long and short
- Rounding of long jobs
- DP for long jobs
- Relaxed decision procedure

Define: job j is long if $p_j > \epsilon T^*$

job j is short if $p_j \leq \epsilon T^*$, for some $T^* \leq OPT$.

PTAS is same as before. Step 2 is easy, and need to find a different algo. for step 1. Choosing T^* is crucial.

Relaxed decision procedure

A simple tool for finding an α -approximation algo. for min problem.

Design an algo. with the following properties:

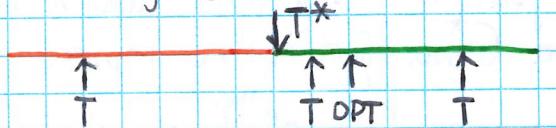
If $T \geq OPT$, it finds a solution of value at most αT

If $T < OPT$

Find the smallest T , say T^* , s.t. algo. finds a solution of value $\leq \alpha T$

Then, we must have $T^* \leq OPT$.

→ Value of solution is at most $\alpha T^* \leq \alpha OPT$



How to find T^* ? (All numbers are integer).

1. Try all values: start with a large value and decrease by 1 each step.
(NOT Polynomial).

2. Binary search (polynomial)

Idea: maintain lower bound LB , and upper bound UB , s.t. at any moment:

i) $OPT \geq LB$

ii) The solution of value $\leq \alpha UB$



Binary search:

while $LB \neq UB$ do:

$$T = L(LB + UB)/2$$

If algo. gives solution of value $\leq \alpha T$, then:

set $UB := T$

else: set $LB := T + 1$

Return T

Running time: $\log(UB - LB)$, which is polynomial in the input size

Remember, we want a schedule for long jobs of length at most $(1+\epsilon)OPT$.

We shall use the relaxed decision procedure and will

design an algo. with following property:

If $T \geq OPT$, it finds a schedule for long jobs of length at most $(1+\epsilon)T$.

We find the smallest T , say T^* , for which the algo. is successful.

Then, $T^* \leq OPT$, and the length of schedule for long jobs is at most $(1+\epsilon)T^* \leq (1+\epsilon)OPT$.

Improving Step 1.

Algo.(T): 1. Say job j is long if $p_j > \epsilon T$.

2. Round down processing times of long jobs:

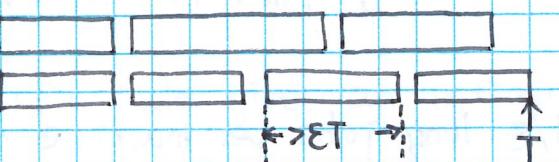
$$p'_j = \lfloor p_j / \mu \rfloor \cdot \mu, \text{ where } \mu = \epsilon^2 / T$$

3. Find a schedule for rounded long jobs of length at most T (by DP)

4. Round up.

Lemma 1: If $T \geq OPT$, then the length of schedule is at most $(1+\epsilon)T$.

Proof: Let σ be the schedule of length at most T computed in step 3.



1) There are at most $T/ET = 1/\epsilon$ long jobs per machine in σ .

2) Rounding up increase the length of each job by at most μ .

1) + 2) → Rounding up increases the total length by at most $(1/\epsilon)\mu = \epsilon T$

Lemma 2: If $T \geq OPT$, then a schedule of length at most T can be found in polynomial time by DP.

Proof: D.P. (Note, if $T \geq OPT$, then such a schedule indeed exists since length are rounded down).

Some definition

• Note: Rounded processing times are multiples of μ .

Say the job is of type t if its rounded processing time is $t\mu$.

Let r be the # of types.

- Let n'_t be the number of long jobs of type t ($t=1, \dots, r$)
- Let (n_1, n_2, \dots, n_r) represent a set of jobs with $n_t \leq n'_t$ jobs of type t ($t=1, \dots, r$)

We call such a vector a configuration

- Let C be the set of all configuration that fit on a single machine i.e. the total job length at most T .

DP-table

For all configuration and # of machines $k \in \{1, \dots, m\}$

Let $F_k(n_1, \dots, n_r) = \text{TRUE}$ iff (n_1, \dots, n_r) can be scheduled on k machines in T .

Filling the DP-table

For all $(n_1, \dots, n_r) \in C$, set $F_1(n_1, \dots, n_r) = \text{TRUE}$

For $k=2 \dots m$:

For all (n_1, \dots, n_r) :

if there is $(s_1, \dots, s_r) \in C$, s.t. $F_{k-1}(n_1 - s_1, \dots, n_r - s_r) = \text{TRUE}$:

Set $F_k(n_1, \dots, n_r) = \text{TRUE}$.

There is a feasible schedule iff $F_m(n'_1, \dots, n'_r) = \text{TRUE}$

Size of DP-table

From Lemma 1. at most $1/\varepsilon$ long jobs per machine

- 1) at most m/ε long jobs in total
- 2) at most m/ε of each type
- 3) at most $(m/\varepsilon)^r$ configurations (n_1, \dots, n_r) (with r the number of types)

- 4) at most $m(m/\varepsilon)^r$ entries $F_k(n_1, \dots, n_r)$ in the DP-table

Rounded processing times are a multiple of μ .

- 5) at most $T/\mu = T/(\varepsilon^2 T) = 1/\varepsilon^2$ different processing times

$$r \leq 1/\varepsilon^2$$

- 4) + 5) \rightarrow DP-table is of polynomial size.

Total running time

- We showed that the number of entries in the DP-table is at most $m(m/\varepsilon)^r$ with $r \leq 1/\varepsilon^2$.
 - Computation of each entry takes $O(|C|)$ time, where, C is the set of all configuration (n_1, \dots, n_r) that fit on a single machine.
 - $|C| \leq (1/\varepsilon)^k$ since there are at most $1/\varepsilon$ long jobs per machine $\rightarrow n_i \leq 1/\varepsilon$
- \Rightarrow Total time is $O(m(m/\varepsilon)^r (1/\varepsilon)^r)$ with $r \leq 1/\varepsilon^2$
which is polynomial in m , assuming that ε is constant.

Chapter 4. Deterministic Rounding of Linear Programs

4.1 Sum of completion times on a single machine.

Let OPT be the optimal value and let OPT^P be the optimal value for the preemptive problem. The optimal non-preemptive schedule is a feasible schedule for the preemptive problem.

$$\text{OPT}^P \leq \text{OPT}.$$

The optimal preemptive schedule is founded by Shortest Remaining Processing Time (SRPT) rule.

Lemma The SRPT rule gives an optimal preemptive schedule.

Proof Assume a schedule σ does not satisfy the SRPT property, then there must be two jobs, say j and k , s.t. the pair j, k does not satisfy the SRPT property.

Let $t = \max\{r_j, r_k\}$. Assume that the remaining processing time of job j at time t is not larger than the remaining processing time of job k at time t .

Now schedule as follows: for every moment $s \geq t$ that σ processes j or k , now process job j at time s until j completes

From that moment only process job k . Let C_j^σ and C_k^σ be the completion time in σ and let C'_j and C'_k be the new completion times. Then,